



 Latest updates: <https://dl.acm.org/doi/10.1145/1328911.1328914>

RESEARCH-ARTICLE

Thin heaps, thick heaps

HAIM KAPLAN, Tel Aviv University, Tel Aviv-Yafo, Tel Aviv District, Israel

ROBERT ENDRE TARJAN, Princeton University, Princeton, NJ, United States

Open Access Support provided by:

[Tel Aviv University](#)

[Princeton University](#)



PDF Download
1328911.1328914.pdf
12 January 2026
Total Citations: 17
Total Downloads: 1008

Published: 28 March 2008
Accepted: 01 September 2007
Revised: 01 September 2006
Received: 01 April 2005

[Citation in BibTeX format](#)

Thin Heaps, Thick Heaps

HAIM KAPLAN

Tel Aviv University

3

AND

ROBERT ENDRE TARJAN

Princeton University and Hewlett Packard Laboratories

Abstract. The Fibonacci heap was devised to provide an especially efficient implementation of Dijkstra's shortest path algorithm. Although asymptotically efficient, it is not as fast in practice as other heap implementations. Expanding on ideas of Høyer [1995], we describe three heap implementations (two versions of *thin heaps* and one of *thick heaps*) that have the same amortized efficiency as Fibonacci heaps, but need less space and promise better practical performance. As part of our development, we fill in a gap in Høyer's analysis.

Categories and Subject Descriptors: F.2 [Analysis of Algorithms and Problem Complexity]; E.2 [Data]: Data Storage Representations

General Terms: Theory, Algorithms

Additional Key Words and Phrases: Data structure, priority queue, heap, fibonacci heap, binomial queue, thin heap, thick heap, melding, decrease key operation

ACM Reference Format:

Kaplan, H. and Tarjan, R. E. 2008. Thin heaps, thick heaps. ACM Trans. Algor. 4, 1, Article 3 (March 2008), 14 pages. DOI = 10.1145/1328911.1328914 <http://doi.acm.org/10.1145/1328911.1328914>

1. Introduction

A *heap* (or *priority queue*) is an abstract data structure consisting of a set of *items*, each with a real-valued *key*, supporting the following operations.

H. Kaplan was supported in part by the Israeli Science Foundation (ISF), Grant No. 548/00. R. E. Tarjan was partially supported at Princeton University by the Aladdin Project, NSF Grant No. CCR-9626862.

Author's addresses: H. Kaplan, School of Computer Science, Tel Aviv University, Israel, e-mail: haimk@math.tau.ac.il; R. E. Tarjan, Department of Computer Science, Princeton University, Princeton, NJ 08544 and Hewlett Packard Laboratories, Palo Alto, CA 94304.

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or direct commercial advantage and that copies show this notice on the first page or initial screen of a display along with the full citation. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, to republish, to post on servers, to redistribute to lists, or to use any component of this work in other works requires prior specific permission and/or a fee. Permissions may be requested from Publications Dept., ACM, Inc., 2 Penn Plaza, Suite 701, New York, NY 10121-0701 USA, fax +1 (212) 869-0481, or permissions@acm.org.
© 2008 ACM 1549-6325/2008/03-ART3 \$5.00 DOI 10.1145/1328911.1328914 <http://doi.acm.org/10.1145/1328911.1328914>

- Create*. Return a new, empty heap.
- Insert*(x, h). Return the heap formed by inserting item x , with predefined key, into heap h .
- Find-min*(h). Return an item in heap h of minimum key.
- Delete-min*(h). Return the heap formed by deleting the item *find-min*(h) from h .
- Meld*(h_1, h_2). Return the heap formed by taking the union of the item-disjoint heaps h_1 and h_2 .
- Decrease-key*(d, x, h). Decrease the key of item x in heap h by subtracting the nonnegative real number d .
- Delete*(x, h). Return the heap formed by deleting item x from heap h .

We assume that any given item is in at most one heap at-a-time and that the operations are allowed to destructively modify their input heaps. In stating bounds, we denote by n the number of items in the heap (or heaps) input to an operation. We assume a random-access machine [Aho et al. 1982] or pointer machine [Ben-Amram 1995; Tarjan 1979] model of computation, so that a memory access takes constant time, and we restrict our attention to data structures that use binary comparisons of keys to make decisions. For heap implementations in a hierarchical memory model (in which memory accesses do not necessarily take constant time), see, for example, various recent papers on so-called *cache-oblivious* data structures [Arge et al. 2002; Brodal and Fagerberg 2002; Brodal et al. 2004]. For heap implementations that use bit-manipulation of keys, see Fredman and Willard [1994], and Thorup [2004]. Finally, we assume that in the case of the operations *decrease-key*(d, x, h) and *delete*(x, h), the position of item x in heap h is known. Without this assumption, if there are arbitrary *meld* operations, a disjoint set data structure must be maintained to keep track of the partition of items into heaps. The paper Kaplan et al. [2002] addresses the effect on the problem of removing this assumption. Finally, we are interested primarily in amortized efficiency [Tarjan 1985], although we make some comments about worst-case efficiency.

Since one can sort by doing n insertions into an initially empty heap followed by n *delete-min* operations, either *insert* or *delete* must have an $\Omega(\log n)$ amortized running time. Many implementations of heaps [Brodal 1996, 1995; Driscoll et al. 1988; Fredman et al. 1986; Fredman and Tarjan 1987; Vuillemin 1978] achieve a worst-case or amortized $O(\log n)$ time bound for all heap operations. One such structure of particular note is the *binomial queue* of Vuillemin [1978], which represents heaps as collections of item-disjoint heap-ordered trees. Binomial queues support all heap operations in $O(\log n)$ worst-case time and are quite efficient in practice [Brown 1978]. Fredman and Tarjan [1987] invented a relaxation of binomial queues called a *Fibonacci heap* that supports *create*, *insert*, *find-min*, *meld*, and *decrease-key* in $O(1)$ amortized time, and *delete-min* and *delete* in $O(\log n)$ amortized time. Fibonacci heaps give an efficient implementation of Dijkstra's algorithm and speed-up algorithms for undirected and directed minimum spanning trees [Fredman and Tarjan 1987; Gabow et al. 1986].

Although theoretically efficient in the amortized sense, Fibonacci heaps are not efficient in the worst case, nor are they as fast or as space efficient in practice as other heap implementations. A straightforward implementation of Fibonacci heaps needs four pointers per node: a parent pointer, left and right sibling pointers, and a first

(leftmost) child pointer. One would prefer an implementation needing only three (or less) pointers per node. More recent work has addressed these issues, among others. Fredman et al. [1986] introduced *pairing heaps*, a self-adjusting variant of Fibonacci heaps. Pairing heaps support all the heap operations in $O(\log n)$ amortized time. They were conjectured to support *decrease-key* in $O(1)$ amortized time, but recent work shows that the amortized time of decrease key is $O(2^{2\sqrt{\log \log n}})$ [Pettie 2005], but $\Omega(\log \log n)$ [Fredman 1999a]. Nevertheless, pairing heaps perform quite well in practice [Stasko and Vitter 1987]. Driscoll et al. [1988] proposed two types of *relaxed heaps*, which, like Fibonacci heaps, are variants of binomial heaps. Both types give a processor-efficient parallel implementation of Dijkstra's shortest path algorithm. *Rank-relaxed heaps* have the same amortized time bounds as Fibonacci heaps. *Run-relaxed heaps* support *create*, *find-min*, and *decrease-key* in $O(1)$ worst-case time, and *delete-min*, *delete*, and *meld* in $O(\log n)$ worst-case time. Brodal [1966], improving on earlier work [Brodal 1995], obtained a (very complicated) heap implementation with the same worst-case time bounds as run-relaxed heaps, but with the worst-case time bound for *meld* improved to $O(1)$. Elmasry improved the bound on comparisons for Fibonacci heaps by a constant factor [Elmasry 2004a] and examined [Elmasry 2004b] versions of pairing heaps, skew heaps [Sleator and Tarjan 1986], and skew-pairing heaps [Fredman 1999b] that use multiway linking instead of binary linking. A *fat heap* [Kaplan and Tarjan 1999] is an earlier data structure that uses a similar idea. Takaoka [2003, 2000] used another form of multiway linking in his 2-3 heaps [Takaoka 2003] and trinomial heaps [Takaoka 2000].

Although the aforementioned structures improve on Fibonacci heaps in various ways, only pairing heaps are as simple conceptually as Fibonacci heaps, but pairing heaps are less efficient theoretically. Our goal in this article is to explore variants of Fibonacci heaps that use less space and are likely to be more efficient in practice, but that retain the theoretical amortized performance of Fibonacci heaps. Interesting work along these lines was done by Peterson [1987] and Høyer [1995]. They both represented a heap by a half-ordered binary tree instead of a heap-ordered forest. The two representations are equivalent; the standard mapping [Knuth 1997] between rooted forests and binary trees maps heap order to half order, as Høyer notes. Peterson's structure is based on half-ordered AVL-trees [Adel'son-Vel'skii and Landis 1962]; Høyer investigates structures related to half-ordered red-black trees [Bayer 1972; Guibas 1978; Tarjan 1982]. The simplest of these structures is Høyer's *one-step heap*.

Our development proceeds independently of Peterson's and Høyer's. We start with the observation that the extra pointer in Fibonacci heaps is the parent pointer, which is needed because of the way *decrease-key* operations work. A *decrease-key* operation can cause a node and its subtree to be cut from its parent, and can further result in a sequence of *cascading cuts*, in which each of a sequence of ancestors is cut from its parent. We consider implementing *decrease-key* in a way that takes advantage of the ordering of siblings in a heap-ordered tree representing a heap: Instead of proceeding from a child to its parent, we proceed from a node to its left sibling, proceeding to its parent only if it has no left sibling. With such a method we need parent pointers only for first children, saving a pointer per node. In Section 2 we describe *thin heaps*, a variant of Fibonacci heaps that uses this idea. An early version of this section appeared in a technical report [Kaplan and Tarjan 1999]. In Section 3 we describe an alternative implementation of *decrease-key* for thin

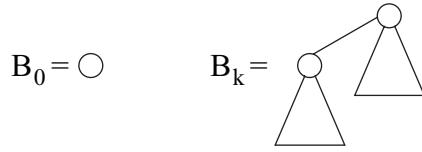


FIG. 1. Binomial trees.

heaps that does only one cut per *decrease-key* (but has an extra subcase). We also describe a related structure, *thick heaps*. Thick heaps turn out to be isomorphic to Høyer's one-step heaps by the heap-ordered forest to half-ordered tree mapping. In the process of our development we fill in a gap in Høyer's analysis; we also find that in the definition of his structure Høyer allows more flexibility than can in fact occur. We conclude in Section 4 with some variants and extensions. We intend our work to provide not only new variants of Høyer's structure, but additional insights and a clear, concise exposition.

2. Thin Heaps

We define binomial trees inductively, as follows. A binomial tree of rank 0 is a single node. For any positive integer k , a binomial tree of rank k consists of two binomial trees of rank $k - 1$ linked so that the root of one becomes the first child of the root of the other (see Figure 1). A binomial tree of rank k contains exactly 2^k nodes, and its root has exactly k children.

Every tree in a thin heap is a binomial tree, each of whose nodes may be missing its first child and the subtree rooted there. More formally, a *thin tree* is an ordered tree, each of whose nodes has a nonnegative integer *rank*, with the following three properties.

- (1) The children of a node with k children have ranks $k - 1, k - 2, \dots, 0$, from first to last.
- (2) A node with k children has rank k or $k + 1$. We call the node *normal* in the former case and *thin* in the latter.
- (3) The root is normal.

The *rank of a thin tree* is the rank of its root. A thin tree all of whose nodes are normal is a binomial tree, and vice-versa. As with binomial trees, if we link two thin trees of equal rank by making the root of one tree the new first child of the root of the other and adding one to the rank of the single remaining root, then the result is a thin tree. Thin trees have the same size bound as the trees in Fibonacci heaps, given by the following lemma, whose proof is more direct than that for Fibonacci heaps.

LEMMA 2.1. *A node with k children in a thin tree has at least $F_{k+1} \geq \phi^k$ descendants, including itself, where F_k is the k th Fibonacci number, defined by $F_0 = 1$, $F_1 = 1$, $F_k = F_{k-2} + F_{k-1}$ for $k \geq 2$, and $\phi = (1 + \sqrt{5})/2$ is the golden ratio.*

PROOF. The inequality $F_{k+1} \geq \phi^k$ is well known [Knuth 1997]. We prove by induction on k that a node with k children in a thin tree has at least F_{k+1} descendants. This is obvious for $k = 0, 1$. Suppose $k \geq 2$, let x be a node with k children in a thin tree, and consider the subtree rooted at x , which is itself a thin tree if we give to x rank k . Cutting the link between x and its first child and decreasing the rank of x by one results in two thin trees: one whose root x has $k - 1$ children and the other whose root has $k - 1$ or $k - 2$ children, by items (1) and (2) of our aforesaid list of properties. By the induction hypothesis, the original number of descendants of x is at least $F_k + \min\{F_k, F_{k-1}\} \geq F_{k+1}$. \square

A *heap-ordered tree* is a rooted tree, each of whose nodes has a real-valued key satisfying *heap order*: No child has key smaller than that of its parent. A *thin heap* is a set of node-disjoint heap-ordered thin trees whose nodes are the heap items. To support the heap operations on thin heaps, we maintain the following data:

- (a) a singly-linked circular list of the tree roots, with a root of minimum key first;
- (b) for each node, a doubly-linked list of its children;
- (c) for each first child, its parent; and
- (d) for each node, its rank.

To maintain this data, an integer (the rank) and three pointers per node suffice: one to the first child; one to the right sibling, or next root if the node is a root; and one to the left sibling, or parent if the node is a first child. This representation supports constant-time linking of two thin trees, given their roots, as well as constant-time tests of whether a node is thin, whether a node is a root, and whether a node is a first child: A node is thin iff its rank is two more than that of its first child, a root iff its left sibling/parent pointer is null, and a first child iff it is the first child of the node indicated by its left sibling/parent pointer. One could also mark nodes as being thin or not; although this is redundant information, storing it explicitly speeds-up testing for thinness.

We perform *create*, *insert*, *find-min*, *meld*, *delete-min*, and *delete* on thin heaps exactly as on Fibonacci heaps. Specifically, to perform *create*, return a null pointer. To perform *insert*(x, h), make a new thin tree with the single node x , and insert x into the root list of h , in second or first position depending on whether its key is larger than that of the old first root. Return the modified heap. To perform *find-min*(h), return the first root of h . To perform *meld*(h_1, h_2), combine the root lists of h_1 and h_2 into a single list whose first root is the first root of h_1 or the first root of h_2 , whichever has smaller key, breaking a tie arbitrarily. Return the new heap.

To perform *delete-min*(h), remove from h the first root of h , say x , and make each child of x normal if it is not by reducing its rank by one. Then combine the list of children of x with the list of roots of h (other than x), and repeat the following *linking step* until it no longer applies: Link two trees of equal rank by making the root of larger key the new first child of the root of smaller key (breaking a tie arbitrarily) and increasing the rank of the new root by one. Once there are no two trees of equal rank, form a list of the remaining roots, choosing a root of minimum key to be the first on the list. Return the modified heap. One can find links to perform in $O(1)$ time per link by using a temporary array indexed by rank to store tree roots; see Fredman and Tarjan [1987]. Alternatively, one can avoid the use of random-access memory and implement *delete-min* (and all the other operations) on

a pointer machine, by having each node of rank k point to a global *rank node* for rank k , and having the rank node for k point to the rank nodes for $k + 1$ and $k - 1$. During a *delete-min* each rank node points to a root of the corresponding rank, or to null.

Perform $\text{delete}(x, h)$ as $\text{delete-min}(\text{decrease-key}(\infty, x, h))$.

The implementation of the last heap operation, *decrease-key*, differs from its implementation on Fibonacci heaps. A *decrease-key* on a Fibonacci heap can cause a sequence of *cascading cuts*, in which successive links along a path of ancestors are cut. A *decrease-key* on a thin tree can cause a sequence of cuts, but only of first children; between such cuts, rank reduction can occur, and the sequence can end with a restructuring step in which a child becomes a sibling of its parent.

Specifically, to perform $\text{decrease-key}(d, x, h)$ on a thin heap, begin by subtracting d from the key of x . If x is a first child whose key remains no less than that of its parent, the *decrease-key* is done. Otherwise, if x is a root, complete the *decrease-key* by making x the first root of h if its key is now smaller than the previous minimum. If, on the other hand, x is not a root, let y be the left sibling of x , or the parent if it is a first child. Cut the link between x and its parent by removing x from the list of children containing it, making it a new tree root. Make x normal if it is not, and add x to the list of roots in first or second position, depending on whether its key is smallest.

Cutting at x may violate properties (1), (2), or (3) at y . Repeat the following *repair step*, which repairs the violation at y but may create a new violation at the left sibling or parent of y , until there is no new violation (see Figure 2). Then return the modified heap.

Repair Step

Case 1. Violation of (1): Node y has rank two greater than that of its next sibling, or has rank 1 and no next sibling.

Case 1a. Node y is thin. Reduce the rank of y by one, repairing the violation and making y normal. Replace y by its left sibling, or by its parent if it is a first child, and check for a violation at the new y .

Case 1b. Node y is normal. Remove the first child of y , say w , and insert w after y in the list of children containing y . This makes node y thin but repairs the violation without creating a new violation.

Case 2. Violation of (2): Node y has rank three greater than that of its first child, or has rank two and no children. Decrease the rank of y by two, repairing the violation and making y normal. Let z be the left sibling of y , or its parent if y is a first child. Remove y from the list of children containing it, and add y to the root list in the second position. Replace y by z and check for a violation at the new y .

Case 3. Violation of (3): Node y is a thin root. Make y normal by decreasing its rank by one.

It is straightforward to verify that *decrease-key* produces a valid thin heap. Note that a node loses at most one child during a *decrease-key*, which means that a root which is initially normal can become thin and violate (3), but can never violate (2). This means that node y in Case 2 is always a nonroot.

The amortized analysis of thin heaps is virtually the same as the analysis of Fibonacci heaps, with thin nodes in thin heaps taking the place of marked nodes

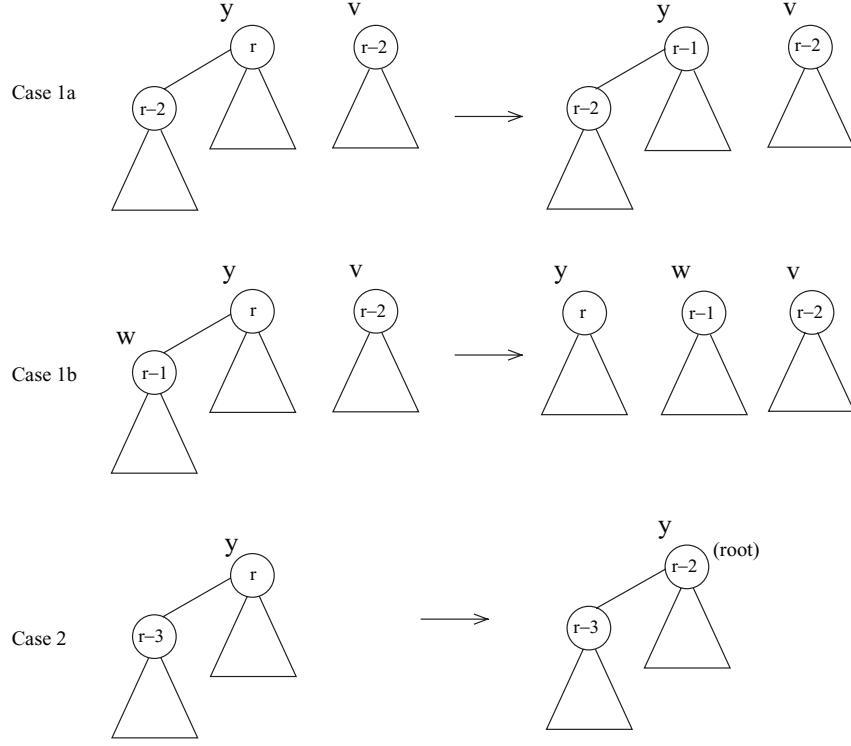


FIG. 2. Nonroot repair step for thin heaps. In Cases 1a and 1b, node v is the right sibling (if it exists) of node y .

in Fibonacci heaps [Fredman and Tarjan 1987]. We use the *potential function* technique [Tarjan 1985]. Assign to each possible collection of node-disjoint thin heaps a nonnegative integer *potential* equal to the number of trees plus twice the number of thin nodes. Define the *amortized time* of a heap operation to be its actual time plus the net increase in potential caused by the operation. Then the total actual time of an arbitrary sequence of heap operations, starting with no heaps, is at most the sum of the amortized times of the operations [Tarjan 1985].

Both the actual and amortized times of *find-min*, *insert*, and *meld* are $O(1)$: An insertion increases the number of trees and hence the potential by one; the other operations do not change the potential. If we charge one unit of time for each linking step, the amortized time of a *delete-min* operation is at most a constant times the maximum node rank, because each linking step reduces the potential by one and the remaining time spent in *delete-min* is bounded by a constant times the maximum rank. By Lemma 2.1, the maximum rank in an n -node thin heap is at most $\log_{\phi} n$, which means that the amortized time of *delete-min* is $O(\log n)$.

To analyze *decrease-key*, let us charge one unit of time for each repair step. Each repair step (except the last) makes a thin node normal and may create a new tree, for a potential drop of one. The initial cut and the last repair step take $O(1)$ time and increase the potential by at most three (at most one for the cut and two for the repair step), resulting in an overall amortized time of $O(1)$ for *decrease-key*. The amortized time of *delete* on an n -node heap is $O(\log n)$, since it is a *decrease-key* followed by a *delete-min*.

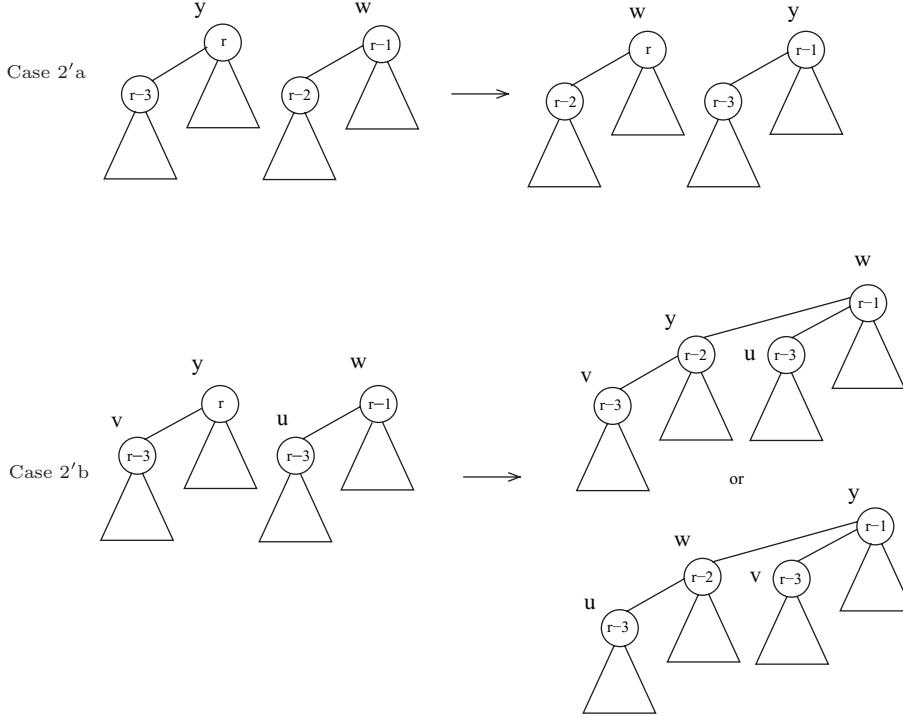


FIG. 3. Alternative repair step for thin heaps. The outcome in Case 2'b depends on which of y and w has smaller key.

Thus we obtain the following theorem.

THEOREM 2.2. *Beginning with no thin heaps and performing an arbitrary sequence of heap operations, if we charge each delete-min and delete on an n -node heap $O(\log n)$ amortized time and each other operation $O(1)$ amortized time, then the total time is at most the total amortized time.*

The worst-case time of a *decrease-key* operation on an n -node thin heap is $O(\log n)$ because each successive violation occurs at a node of higher rank. In contrast, a single *decrease-key* on an n -node Fibonacci heap can take $\Omega(n)$ time, as one can show by modifying the solution to exercise 20.4-1 on page 496 of Cormen et al. [2001].

3. From Thin Heaps to Thick Heaps

As on a Fibonacci heap, a *decrease-key* operation on a thin heap can in the worst case take more than one cut. By modifying Case 2 of *decrease-key*, however, we can avoid all but the first cut, the one at the node whose key decreases. Specifically, we replace Case 2 by the following variant (see Figure 3).

Case 2'. Violation of (2): Node y has rank three greater than that of its first child, or has rank two and no children. Let w be the right sibling of node y .

Case 2'a. Node w is normal. Increase the rank of w by one, decrease the rank of y by one, and swap y and w in the list of children containing them. This makes both y and w thin and repairs the violation without creating a new violation.

Case 2'b. Node w is thin. Let z be the left sibling of y , or its parent if y has no left sibling. If y has key no less than that of w , remove y from the list of children containing it and make it the new first child of w ; decrease the rank of y by two. If, on the other hand, y has key less than that of w , remove w from the list of children and make it the new first child of y ; decrease the ranks of both w and y by one. (In either case, both y and w become normal.) Finally, replace y by z and check for a violation at the new y .

As with the original implementation, it is straightforward to verify that this implementation of *decrease-key* produces a valid thin heap. The amortized analysis is similar, except that a smaller potential function suffices. Specifically, we define the potential of a collection of thin heaps to be the number of trees plus the number of thin nodes. The only nonterminating cases in *decrease-key*, Case 1a and Case 2'b, reduce the potential by one and two, respectively. The rest of the analysis is the same as in Section 2. Thus, we obtain Theorem 2.2 for thin heaps with the alternative *decrease-key* implementation.

We can also try to improve thin heaps by reducing the number of children of each node. We can in fact do this by replacing property (2) with the following alternative.

(2') A node with k children has rank k or $k - 1$. We call the node *normal* in the first case and *thick* in the latter.

We call a tree with node ranks that satisfy (1), (2'), and (3) a *thick tree*. For thick trees the following improvement of Lemma 2.1 holds.

LEMMA 3.1. *A node with k children in a thick tree has at least 2^k descendants, including itself.*

PROOF. We prove the lemma by induction on k . The lemma is clearly true for $k = 0$. Suppose $k \geq 1$, let x be a node with k children in a thick tree, and consider the subtree rooted at x , which is itself a thick tree if we give to x rank k . Cutting the link between x and its first child and decreasing the rank of x by one results in two thick trees: one whose root x has $k - 1$ children and the other whose root, say y , has k or $k - 1$ children. If y has k children, we obtain a smaller thick tree whose root has $k - 1$ children by cutting the link between y and its first child. Thus, the original subtree rooted at x consists of at least two disjoint trees with roots having $k - 1$ children, plus possibly additional nodes. By the induction hypothesis the original subtree rooted at x contains at least $2^{k-1} + 2^{k-1} = 2^k$ nodes. \square

A *thick heap* is a set of node-disjoint heap-ordered thick trees whose nodes are the heap items. We can implement all the heap operations except *decrease-key* on thick heaps exactly as on thin heaps. The implementation of *decrease-key* is similar to the alternative implementation for thin heaps, with the analogs of Cases 1a and 1b (as well as those of Cases 2'a and 2'b) switched. Specifically, perform *decrease-key* on a thick heap exactly as on a thin heap, but use the following repair step (see Figure 4).

Repair Step for Thick Heaps

Case 1''. Violation of (1): Node y has rank two greater than that of its next sibling, or has rank 1 and no next sibling.

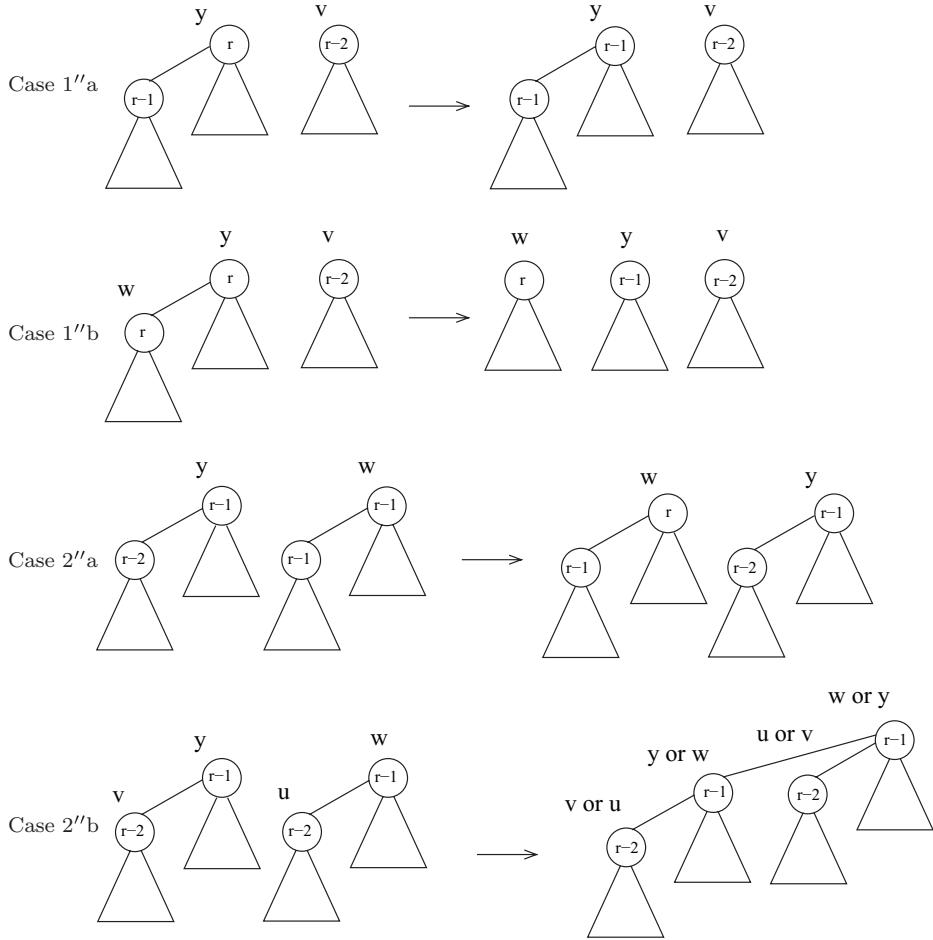


FIG. 4. Repair step for thick heaps. In Cases 2''a and 2''b, node y is a nonroot and its rank has already been decreased by one. The outcome in Case 2''b depends on which of y and w has smaller key.

Case 1''a. Node y is normal. Reduce the rank of y by one, repairing the violation and making y thick. Replace y by its left sibling, or by its parent if it is a first child, and check for a violation at the new y .

Case 1''b. Node y is thick. Remove the first child of y , say w , decrease the rank of y by one, and insert w before y in the list of children containing y . This leaves y thick and repairs the violation without creating a new violation.

Case 2''. Violation of (2'): Node y has rank two greater than that of its first child, or has rank one and no children. Decrease the rank of y by one, making it normal. If y is a root, the repair is complete. Otherwise, let w be the right sibling of y and apply the appropriate one of the following two cases.

Case 2''a. Node w is thick. Increase the rank of w by one and swap y and w in the list of children containing them. This makes both y and w normal and repairs the violation without creating a new violation.

Case 2''b. Node w is normal. Let z be the left sibling of y , or its parent if it has no left sibling. If y has key no less than that of w , remove y from the list of children

containing it and make it a new first child of w . (Node w becomes thick.) If, on the other hand, y has key less than that of w , remove w from the list of children containing it and make it a new first child of y . (Node y becomes thick.) Finally, replace y by z and check for a violation at the new y .

The correctness of *decrease-key* is easy to establish. Note that a violation of (3) can occur only if the violating root violates (2') as well. Thus, the analog of Case (3) is part of Case 2''. More interesting is the amortized analysis of thick heaps. In thin heaps, thin nodes are bad in that they can cause expensive *decrease-key* operations. In thick heaps, normal (nonroot) nodes are bad in the same way; thick nodes are good. Our potential function for thick heaps is twice the number of trees plus the number of nonroot normal nodes. With this definition of the potential, the amortized time of a link (in a *delete-min*) is zero if we charge it an actual time of one; the link reduces the number of trees by one (decreases the potential by two) but increases the number of nonroot normal nodes by one (increases the potential by one). Each of the nonterminating cases in *decrease-key* (1''a and 2''b) converts a previously normal node to thick, decreasing the potential by one; hence, the amortized time to apply the case is zero if we charge it an actual time of one. The remainder of the analysis is exactly as for thin heaps, and thus we obtain Theorem 2.2 for thick heaps. As with thin heaps, the worst-case time of a *decrease-key* on an n -node thick heap is $O(\log n)$ because each successive repair step is on a node of higher rank.

As noted in Section 1, thick heaps are isomorphic to Høyer's one-step heaps by the heap-ordered-forest to half-ordered-tree mapping. In addition to this representational difference, our presentation differs from Høyer's in three ways. First, Høyer relaxes property (2') to allow a node to have any rank no greater than its number of children. In fact, all the operations create only normal and thick nodes; Høyer's relaxation adds no generality. Second, Høyer adds to his data structure what he calls a *tree list*, which has the effect of reducing the worst-case time for a *delete-min* from $\Theta(n)$ to $O(\log n)$, while increasing the worst-case time for a *meld* from $O(1)$ to $\Theta(\log n)$. The tree list is a variant of an idea introduced by Driscoll et al. [1988] and is related to the idea of a redundant counter, used later by Brodal [1996, 1995] and Kaplan et al. [2002]. We discuss the use of a tree list briefly in Section 4. Finally, Høyer's amortized analysis of one-step heaps has a lacuna; he does not account for the creation of nonroot normal nodes by links. The preceding analysis fills this gap.

The implementation of *decrease-key* on thick heaps is analogous to the alternative implementation on thin heaps. It is natural to consider an implementation analogous to the original one on thin heaps. We obtain such an implementation by making the following substitution for Case 2''.

Case 2'''. Violation of (2'): Node y has rank two greater than that of its first child, or has rank two and no children. Decrease the rank of y by one, making y normal. If y is not a root, proceed exactly as in Case 2, removing y from its list of children, adding it to the root list, and checking for a new violation at the old left sibling of y , or at its old parent if it had no left sibling.

Perhaps surprisingly, thick heaps with this implementation of *decrease-key* are asymptotically less efficient. To obtain an expensive sequence of operations, do $2^k + 1$ insertions followed by one *delete-min*, resulting in a thick heap that is a B_k tree with all nodes normal. Then repeat the following sequence of operations any

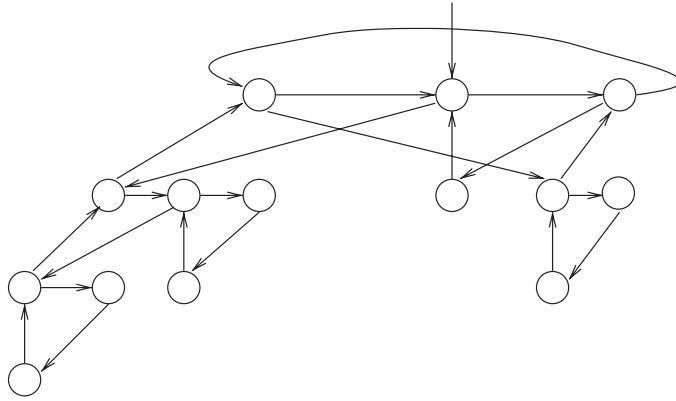


FIG. 5. Two-pointer-per-node representation of a heap. The heap consists of three binomial trees: a B_3 , a B_1 , and a B_2 . This representation is essentially the reverse of Brown's "structure K."

number of times: Do k *decrease-key* operations on the deepest nodes of rank zero, followed by a single insert and a single *delete-min*. This sequence takes $\Theta(k^2)$ time and reproduces a B_k tree with all nodes normal. The amortized cost per *decrease-key* is $\Theta(k) = \Theta(\log n)$.

4. Variants and Extensions

All the variants and extensions discussed here apply equally to both versions of thin heaps and to thick heaps, and also to other heap structures presented by Høyer and similar structures as well.

We can reduce the number of pointers per node from three to two at the cost of a small constant factor in running time. Brown [1978, pp. 306–308] describes how to do this for binomial queues; the same idea works for thin or thick heaps. For example, each node can point to its right sibling, or to the next root if it is a root; and to its parent if it is a first child, or to the first child of its left sibling if it is not a root and not a first child, or to the first child of its previous root if it is a root (see Figure 5). With this representation, there is no easy way to test whether a node is a root. To deal with this, we use an extra bit to mark the roots.

Gabow et al. [1986] gave an application of heaps to the problem of finding minimum spanning trees in directed graphs. It requires an additional heap operation, that of moving a node x , along with all of its descendants, from one heap to another. This operation is only allowed if x does not have minimum key in its heap. Gabow et al. [1986] describe how to perform this operation on a Fibonacci heap in $O(1)$ amortized time; their method also works on thin or thick heaps. A move is just like a *decrease-key* operation, except that the tree with root x is moved to the new heap.

On thin and thick heaps the worst-case time of *delete-min* is $O(n)$. We can reduce this to $O(\log n)$ without affecting the amortized time bounds of any of the operations, as follows. Keep track of the number of trees and the number of nodes in each heap. Whenever the number of trees in an n -node heap exceeds $2 \log_\phi n$ for thin heaps or $2 \log n$ for thick heaps, do a cleanup by repeatedly linking trees of equal rank until there is at most one tree per rank. Such a cleanup takes zero amortized time because the $O(\log n)$ actual time is paid for by the drop in potential caused by the

links. If cleanups are done, the amortized time of all the heap operations remains the same; the worst-case times are $O(\log n)$ for all operations except *find-min*, which remains $O(1)$.

A refinement of the cleanup idea preserves the $O(\log n)$ worst-case time for *delete-min*, *meld*, and *decrease-key*, while reducing the worst-case time for *insert* back to $O(1)$. This is Høyer's tree list method: Maintain, for each heap, a list of pairs of trees of equal rank, plus at most one “odd” tree per rank. When adding a new tree by an *insert* or a *decrease-key*, link two paired trees, if there are any. This guarantees that the number of trees is $O(\log n)$. To obtain an $O(1)$ amortized time bound for *meld*, one must add $O(\log n)$ to the potential of an n -node heap; this pays for combining the tree lists during the *meld* of two heaps. A drawback of this method is that each heap requires an array of trees, one per rank; there is no obvious way to implement this idea on a pointer machine.

REFERENCES

- ADEL'SON-VEL'SKII, G. M., AND LANDIS, E. M. 1962. An algorithm for organization of information. *Dokl. Akad. Nauk SSSR* 146, 263–266.
- AHO, A. V., HOPCROFT, J. E., AND ULLMAN, J. D. 1982. *Data Structures and Algorithms*. Addison-Wesley.
- ARGE, L., BENDER, M. A., DEMAINE, E. D., HOLLAND-MINKLEY, B., AND MUNRO, J. I. 2002. Cache-Oblivious priority queue and graph algorithm applications. In *Proceedings of the 34th ACM Symposium on Theory of Computing (STOC)*. ACM Press, 268–276.
- BAYER, R. 1972. Symmetric binary b-trees: Data structure and maintenance algorithms. *Acta Inf.* 1, 290–306.
- BEN-AMRAM, A. M. 1995. What is a “pointer machine”? *SIGACT News* 26, 2, 88–95.
- BRODAL, G. S. 1996. Worst-Case priority queues. In *Proceedings of the 7th Annual ACM-SIAM Symposium on Discrete Algorithms (SODA)*. ACM Press, 52–58.
- BRODAL, G. S. 1995. Fast meldable priority queues. In *Proceedings of the 4th International Workshop on Algorithms and Data Structures (WADS)*. Lecture Notes in Computer Science, vol. 955. Springer, 282–290.
- BRODAL, G. S., AND FAGERBERG, R. 2002. Funnel heap—A cache oblivious priority queue. In *Proceedings of the 13th Annual International Symposium on Algorithms and Computation (ICALP)*. Lecture Notes in Computer Science, vol. 2518. Springer, 219–228.
- BRODAL, G. S., FAGERBERG, R., MEYER, U., AND ZEH, N. 2004. Cache-Oblivious data structures and algorithms for undirected breadth-first search and shortest paths. In *Proceedings of the 9th Scandinavian Workshop on Algorithm Theory (SWAT)*. Lecture Notes in Computer Science, vol. 3111. Springer, 480–492.
- BROWN, M. R. 1978. Implementation and analysis of binomial queue algorithms. *SIAM J. Comput.* 7, 3, 298–319.
- CORMEN, T. H., LEISERSON, C. E., RIVEST, R. L., AND STEIN, C. 2001. *Introduction to Algorithms*, 2nd ed. MIT Press, Cambridge, MA.
- DRISCOLL, J. R., GABOW, H. N., SHRAIRMAN, R., AND TARJAN, R. E. 1988. Relaxed heaps: An alternative to Fibonacci heaps with applications to parallel computation. *Commun. ACM* 31, 11, 1343–1354.
- ELMASRY, A. 2004a. Layered heaps. In *Proceedings of the 9th Scandinavian Workshop on Algorithm Theory (SWAT)*. Lecture Notes in Computer Science, vol. 3111. Springer.
- ELMASRY, A. 2004b. Parameterized self-adjusting heaps. *J. Algor.* 52, 2, 103–119.
- FREDMAN, M. L. 1999a. On the efficiency of pairing heaps and related data structures. *J. ACM* 46, 4, 473–501.
- FREDMAN, M. L. 1999b. A priority queue transform. In *Proceedings of the 3rd International Workshop on Algorithm Engineering (WAE)*. Lecture Notes in Computer Science, vol. 1668. Springer, 243–257.
- FREDMAN, M. L., AND WILLARD, D. E. 1994. Trans-Dichotomous algorithms for minimum spanning trees and shortest paths. *J. Comput. Syst. Sci.* 48, 3, 533–551.
- FREDMAN, M. L., AND TARJAN, R. E. 1987. Fibonacci heaps and their uses in improved network optimization algorithms. *J. ACM* 34, 3, 596–615.

- FREDMAN, M. L., SEDGEWICK, R., SLEATOR, D. D., AND TARJAN, R. E. 1986. The pairing heap: A new form of self-adjusting heap. *Algorithmica* 1, 111–129.
- GABOW, H. N., GALIL, Z., SPENCER, T., AND TARJAN, R. E. 1986. Efficient algorithms for finding minimum spanning trees in directed and undirected graphs. *Combinatorica* 6, 2, 109–122.
- GUBAS, L. J., AND SEDGEWICK, R. 1978. A dichromatic framework for balanced trees. In *Proceedings of the Annual Symposium on Foundations of Computer Science (FOCS)*, 8–21.
- HOYER, P. 1995. A general technique for implementation of efficient priority queues. In *Proceedings of the 3rd Israel Symposium on Theory of Computer Systems*. IEEE Computer Society Press, 57–66.
- KAPLAN, H., AND TARJAN, R. E. 1999. New heap data structures. Tech. Rep. TR-597-99, Princeton University.
- KAPLAN, H., SHAFRIR, N., AND TARJAN, R. E. 2002. Meldable heaps and Boolean union-find. In *Proceedings of the 34th Annual ACM Symposium on Theory of Computing (STOC)*, 573–582.
- KNUTH, D. E. 1997. *The Art of Computer Programming*, Volume 1: *Fundamental Algorithms*, 3rd ed. Addison-Wesley, Reading, MA.
- PETERSON, G. L. 1987. A balanced tree scheme for meldable heaps with updates. Tech. Rep. GIT-ICS-87-23, School of Information and Computer Science, Georgia Institute of Technology, Atlanta, Georgia.
- PETTIE, S. 2005. Towards a final analysis of pairing heaps. In *Proceedings of the 46th Annual Symposium on Foundations of Computer Science (FOCS)*, 174–183.
- SLEATOR, D., AND TARJAN, R. E. 1986. Self-Adjusting heaps. *SIAM J. Comput.* 15, 1, 52–69.
- STASKO, J. T., AND VITTER, J. S. 1987. Pairing heaps: Experiments and analysis. *Commun. ACM* 30, 3, 234–249.
- TAKAOKA, T. 2003. Theory of 2-3 heaps. *Discrete Appl. Math.* 126, 1, 115–128.
- TAKAOKA, T. 2000. Theory of trinomial heaps. In *Proceedings of the 6th Annual International Conference on Computing and Combinatorics (COCOON)*. Lecture Notes in Computer Science, vol. 1858. Springer, 362–372.
- TARJAN, R. E. 1979. A class of algorithms which require nonlinear time to maintain disjoint sets. *J. Comput. Syst. Sci.* 18, 2, 110–127.
- TARJAN, R. E. 1985. Amortized computational complexity. *SIAM J. Algebr. Discrete Meth.* 6, 2, 306–318.
- TARJAN, R. E. 1982. *Data Structures and Network Algorithms*. SIAM, Philadelphia, PA.
- THORUP, M. 2004. Integer priority queues with decrease key in constant time and the single source shortest paths problem. *J. Comput. Syst. Sci.* 69, 3, 330–353.
- VUILLEMIN, J. 1978. A data structure for manipulating priority queues. *Commun. ACM* 21, 309–314.

RECEIVED APRIL 2005; REVISED SEPTEMBER 2006; ACCEPTED SEPTEMBER 2007